

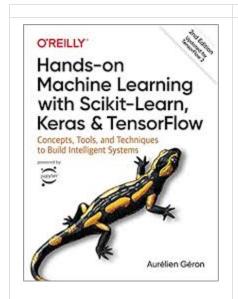
Overview

- one same project of object oriented design and implementation in python
- project is a machine learning method, random forests
- applied to 3 important tasks in machine learning: classification, regression, anomaly detection, plus
 - different impurity measures
 - optimized version "extra trees"
 - feature importance
 - Python coding style and logging
- more importantly, all this needs to apply object oriented principles and several design patterns

Why object-oriented

- It *is* possible to build a random forest classifier without classes, just a bunch of functions, like in this messy implementation.
- But then *extending* it to regression, anomaly detection, compute feature importance etc. becomes difficult.
- By applying OO principles and patterns we can build a design and implementation where such extensions, and probably others, can be added gracefully, easily.
- At the same time following a code style will make the code much more easy to understand.

Reference



Hands-on machine learning with Scikit-Learn, Keras and TensorFlow: concepts, tools, and techniques to build intelligent systems. Aurélien Géron, O'Reilly (2019)

Electronic version at UAB library

What we need is in pages 1-31 (Ch. 1), 175-186 (Ch. 6) and 189-199 (Ch. 7). This selected pages are also in campus virtual.

You can read also the Wikipedia entry.

Marks

- the project is divided into 3 steps or milestones
- what to do exactly for each milestone is in a set of slides I'll show you along the course
- then the goals of the next slides will become clear
- each milestone has its own grading scheme, see next slides
- the possible marks of a milestone are A=10, B=7.5, C=5, D=2.5, E=0.
- you get E when no show, the work is a copy or does not achive the minimum set for D

First milestone

Goals

- design and implement a random forest for classification
- design in UML notation
- test it on two small datasets: Sonar and Iris
- start with the Gini Index as the only split criterion and then extend to others
- optimize the algorithm in two ways: "extra-trees" and parallelize it
- test the optimized versions on the larger MNIST dataset
- practice logging in Python
- start choosing good names for classes, methods etc. and write good comments

Grading

- C:
 - i. the design follows OO principles
 - ii. only the Gini index
 - iii. code works well on the two small datasets
 - iv. the UML design has enough detail and is right and correspond to the code
 - v. logging sentences to debug/trace the execution instead of print
 - vi. right names for classes, methods, attributes, variables, constants + comment for each class and key methods

- B : C plus
 - i. design easily extendable with new impurity measures by applying a certain design pattern
 - ii. added entropy
- A : B plus
 - i. optimization by parallelism, small number of values and extra-trees
 - ii. comparison of elapsed times, both for training and prediction
 - iii. good accuracy on MNIST for each optimization
- D : the classifier works fairly well and code corresponds to an existing UML design, some but not all of the necessary OO principles have been applied

Deliverables

Upload into campus virtual a file named first_milestone.zip with

- 1. a text file autors.txt with the NIU and names of authors
- 2. a file source_code.zip with Python source code, the plantUML class diagram design.puml and a design.png image of the UML design (exported)
- 3. a file results.pdf (export a presentation) with the parameters of the experiments for Sonar and Iris and the accuracies obtained
- 4. if you have done the optimization part, a file times.pdf (export a presentation) with the tables comparing exection times

Do not include the datasets.

Second milestone

Goals

- modify and extend the design to also do regression, avoiding duplicate code by employing inheritance and a certain desing pattern
- print the learned decision trees and compute feature importance by applying another design pattern
- keep the good practices of logging, comments and naming

Grading

• C:

- i. regression part is done and works
- ii. the test with the minimum daily temperature dataset achieves low RMSE error
- iii. the UML class design is updated with regression
- iv. more logging sentences to debug/trace the execution
- v. right names for classes, methods etc. plus comments

• B : C plus

- i. feature importance and print trees done by applying a certain design pattern
- ii. the UML class design is updated with feature importance
- iii. tree printing and feature importance work well on small datasets Sonar and Iris

- A : B plus
 - feature importance on MNIST using extra-trees produces a reasonable result
- D : C.i is done but one or more of the other items in C are wrong or missing

Deliverables

Within a file second_milestone.zip put items 1 and 2 like in the first milestone, plus a file results.pdf with:

- figures of regression and RMSE achieved for the prediction of the last 1 and 2 years of the dataset min. daily temperature
- two printed trees for Iris
- computed importances of features of datasets Sonar and Iris
- figure of the importance of each pixel for the dataset MNIST, shown as an image, adding as scale a plt.colorbar()

Third milestone

Goals

- extend the design to accomodate a new machine learning task, anomaly (outlier) detection
- keep doing logging
- edit the code to conform to Python coding style PEP-8
- write good comments and documenting
- use assertions to check some methods contract

Grading

- C:
 - i. anomaly detection done with the isolation forest algorithm works well on the synthetic and credit card fraud datasets
 - ii. the UML design has been updated accordingly
 - iii. comments are at the right places and with the right content
- B : C plus
 - i. the code has no PEP-8 problems or warnings, or very few, and when they exists there is a reason explained in a comment
 - ii. the classes intended to the users have been documented with doc strings (avoiding copy-paste), and their documentation generated with Sphinx

• A : B plus

- isolation forest, which requires the extra-trees optimization, finds reasonable anomalies in the MNIST dataset, and precision and recall are similar to those reported in the slides
- type hints for every function, class method, argument and class attribute
- D: isolation forest design and implementation is not perfect but still detect some ouliers, or the UML design is wrong or missing

Deliverables

Within a file third_milestone.zip put items 1 and 2 like in the first milestone, plus

- a file results.pdf with:
 - result figure for synthetic dataset
 - what are your outliers within the first 10K samples of the credit card fraud dataset and top 0.5% scores, the precision and recall
 - outliers detected for the MNIST dataset for 3 digits at least, for top 0.2% scores
- a file style.txt with the remaining PEP-8 problems and why
- the documentation generated by Sphinx